Fundamentals of Parallel Systems Lab Book

Jordan Stephano Gray

40087220

graybostephano@gmail.com

April 2016

1   Exercise 2.1

The Producer takes in an integer based on user input. If the integer is less than or equal to 0 the program finishes. The program will only accept integers less than or equal to 100. The result of the input is written out to the channel.

Producer

Out

Connection1

In

The Multiplier has a variable called factor which is defined as 2. An integer with the label "i" is then assigned whichever value was passed in through the previous channel and multiplies that value by the factor (2). This new value is then written to the next channel. The next value of i is then read.

Multiplier

Out

Connection2

In

This does the same as the previous multiplier which results in a value that has been multiplied by a factor of 2, twice. The result is then passed out to the next channel.

Multiplier

Out

Connection3

In

This takes in the value from the previous channel and sets it to a variable "i". The result is then printed in the console and the next "i" value is read.
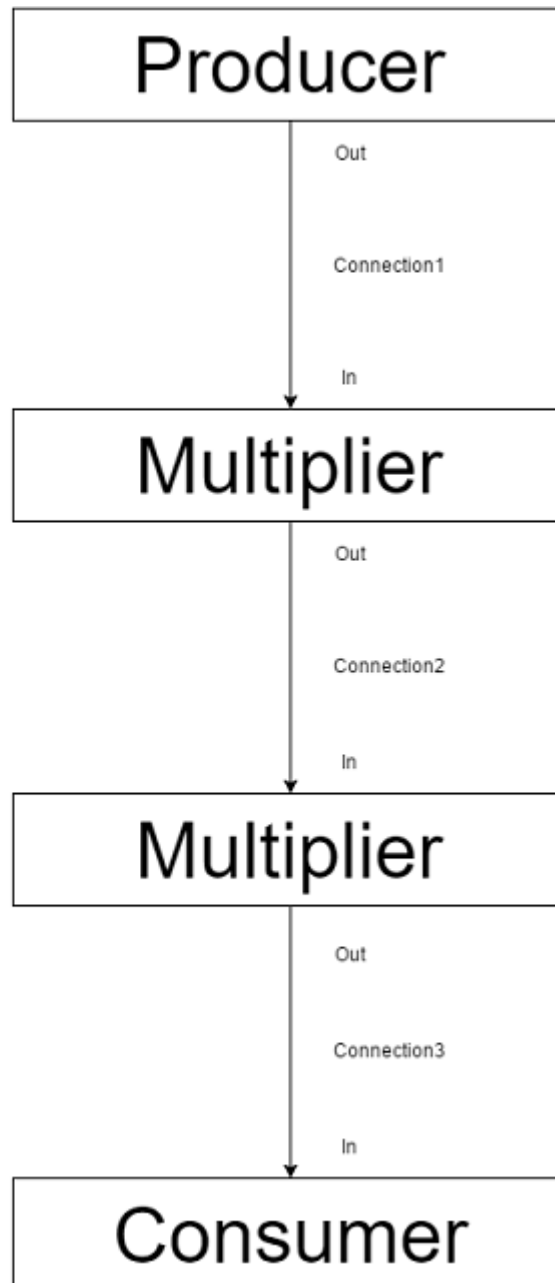
Consumer

Figure 1: Diagram of the system for multiplication.

1.1 Consumer

```
void run() {
  def i = inChannel.read()
  while ( i > 0 ) {
    //Insert modified print line statement
    println "\nResult: $i"
    i = inChannel.read()
  }
  println "Finished"
}
```
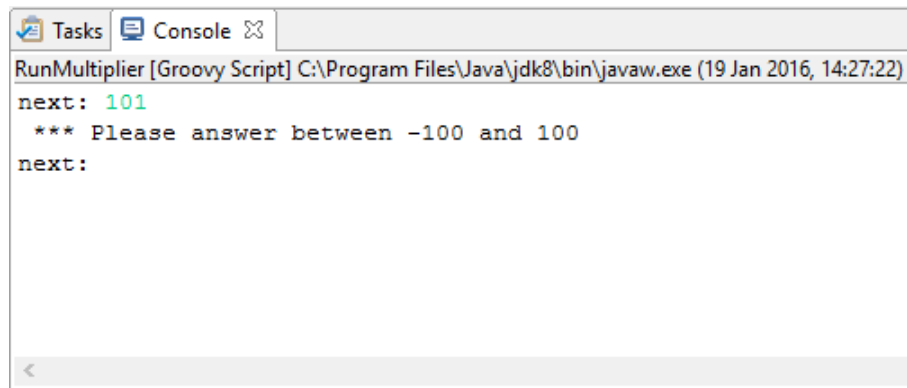
1.2 Multiplier

```
class Multiplier implements CSProcess {

  def ChannelOutput outChannel
  def ChannelInput inChannel
  def int factor = 2

  void run() {
    def i = inChannel.read()
    while (i > 0) {
      // write i * factor to outChannel
      outChannel.write( i * factor )
      // read in the next value of i
      i = inChannel.read()

    }
    outChannel.write(i)
  }
}
```
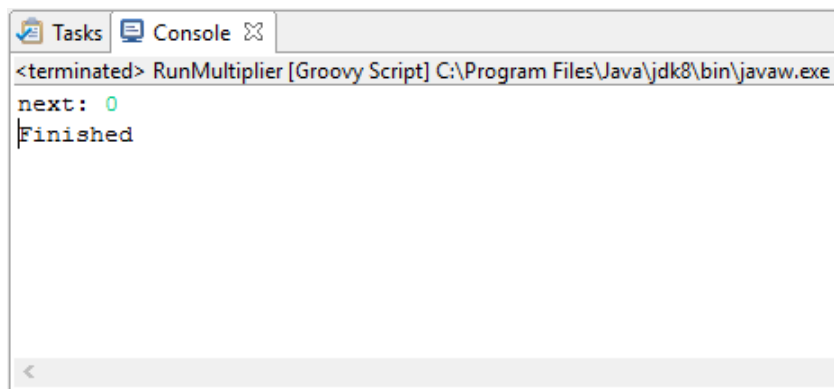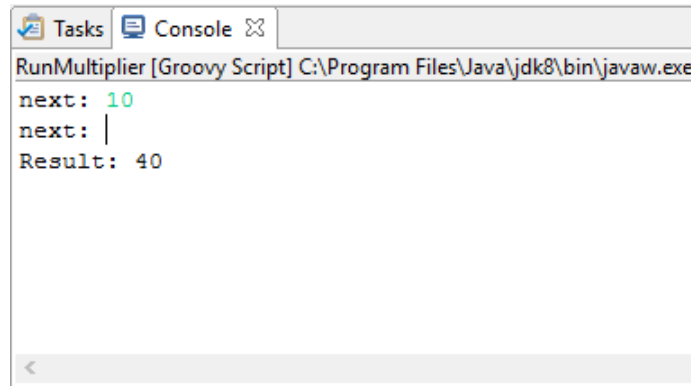
Figure 2: This demonstrates what happens when an input outside of the range (-100 - 100) is used.

Figure 3: This demonstrates what happens when 0 is used as an input and the program terminates.

Figure 4: This demonstrates what happens when a valid input is used.

2 Exercise 2.2

Q: What change is required to output objects containing six integers?

Ans: Modify the for loop used in "CreateSetsOfEight" to loop for every index from 0 to 5 rather than 0 to 7.

Q: How could you parametrize this in the system to output objects that contain any number of integers?

Ans: Create an integer variable to hold the list parameter and substitute it into the for loop in "CreateSetsOfEight". The variable is then easily changed in "RunThreeToEight".

Q: What happens if number of integers required in the output stream is not a factor of total number of integers in input stream?

Ans: Any remaining numbers from input stream are added to an object but because the object is never full it is left out of the output stream and the program can't terminate.

GenerateSetsOfThree defines a list with each index holding three numbers and writes this list to the outChannel.
ListToStream then takes this data through the inChannel and separates these groups of three into individual integers and streams them to the outChannel.
GenerateSetsOfEight then defines a new list to hold the new assortment of values, the listParameter is set to 7 which means the data will be split into sets of 8 (including index 0 – 7). After the values have been amended to the output list the list is printed of with a simple print statement.
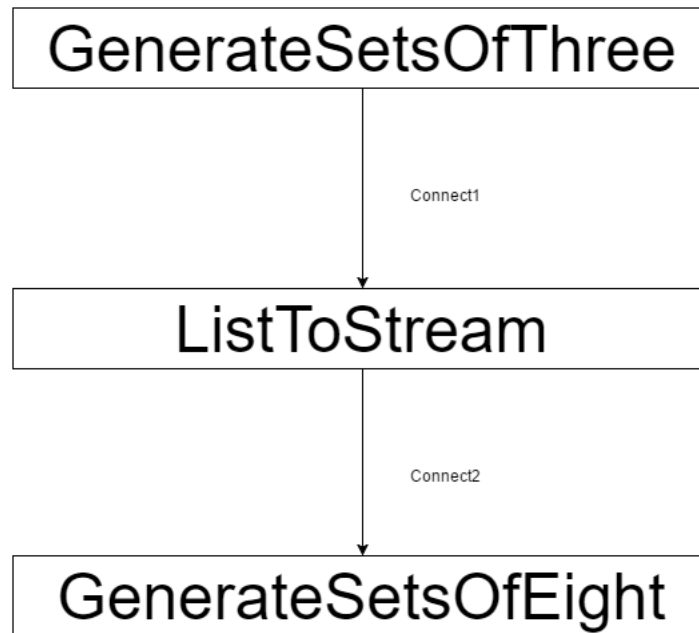
5

Figure 5: Diagram showing the structure of exercise 2-2.

2.1 ListToStream

```
void run (){
   def inList = inChannel.read()
   while (inList[0] != -1) {
      // hint: output list elements as single integers
      outChannel.write(inList[0])
      outChannel.write(inList[1])
      outChannel.write(inList[2])
      inList = inChannel.read()
   }
   outChannel.write(-1)
}
```
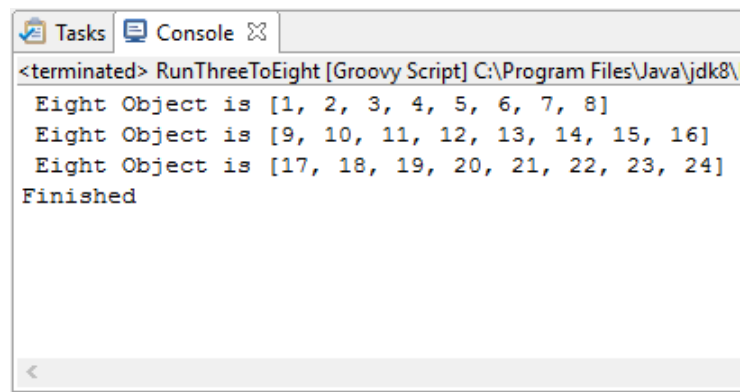
2.2 CreateSetsOfEight

```
def ChannelInput inChannel
def int listParam = 7

void run(){
```

```
    def outList = []
    def v = inChannel.read()
    while (v != -1){
      for ( i in 0 .. listParam ) {
         // put v into outList and read next input
         outList << v
         v = inChannel.read()
      }
      println " Eight Object is ${outList}"
      outList = []
    }
    println "Finished"
}
```
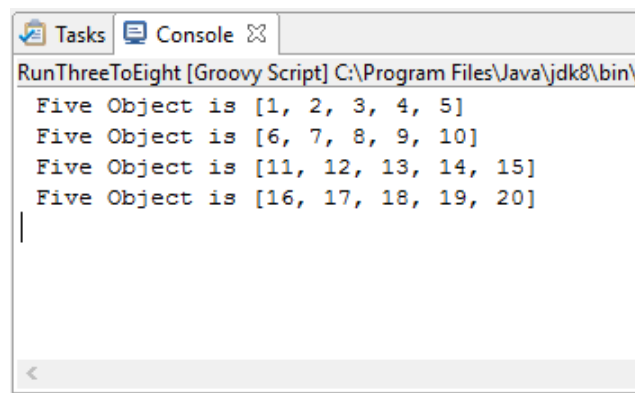
Eight Object is [1, 2, 3, 4, 5, 6, 7, 8]
Eight Object is [9, 10, 11, 12, 13, 14, 15, 16]
Eight Object is [17, 18, 19, 20, 21, 22, 23, 24]
Finished

Figure 6: Shows the program producing eights.

Five Object is [1, 2, 3, 4, 5]
Five Object is [6, 7, 8, 9, 10]
Five Object is [11, 12, 13, 14, 15]
Five Object is [16, 17, 18, 19, 20]

Figure 7: Shows the program producing fives with left overs.

3 Exercise 3.1

Q: Which is the more pleasing solution and why?

Ans: I found that using the Minus method was more pleasing to me. Although using the Negator makes more sense by just inverting the value to a negative number although there is also an extra intermediate step.
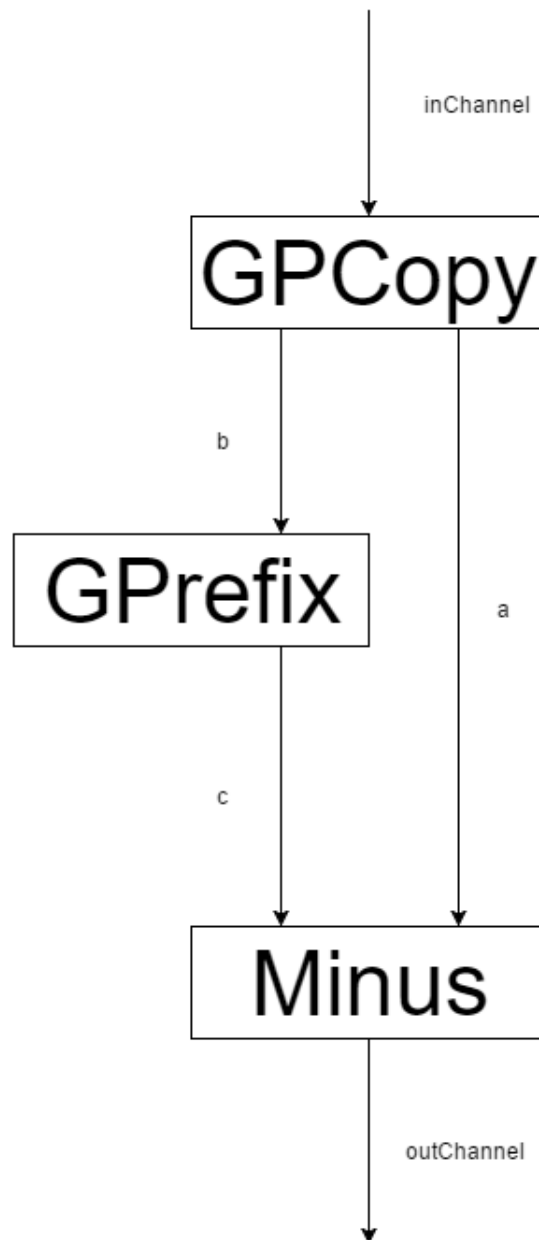
inChannel

GPCopy

b

GPrefix

a

c

Minus

outChannel

Figure 8: System using the Minus function.

9
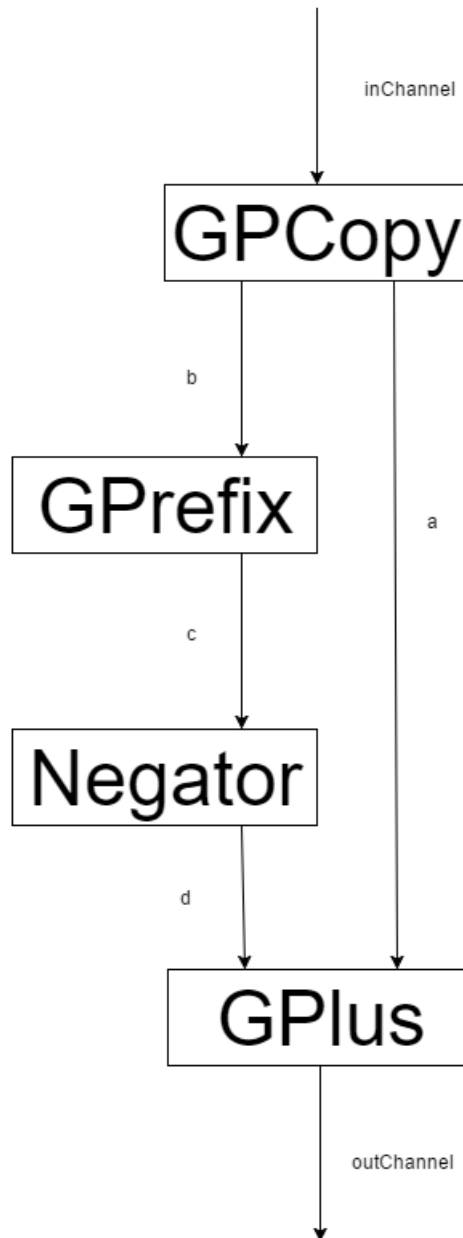
Figure 9: System using the Negator function.

3.1 Differentiate

```
   def differentiateList = [ new GPrefix ( prefixValue: 0,
                                           inChannel: b.in(),
                                           outChannel: c.out() ),
                             new GPCopy ( inChannel: inChannel,
                                          outChannel0: a.out(),
                                          outChannel1: b.out() ),
                             // insert a constructor for Minus
                       new Minus ( inChannel0: a.in(),
                                   inChannel1: c.in(),
                                   outChannel: outChannel )
                   ]

   new PAR ( differentiateList ).run()

 }
```

3.2 DifferentiateNeg

```
   def differentiateList = [ new GPrefix ( prefixValue: 0,
                                           inChannel: b.in(),
                                           outChannel: c.out() ),
                             new GPCopy ( inChannel: inChannel,
                                          outChannel0: a.out(),
                                          outChannel1: b.out() ),
                             //insert a constructor for Negator
                       new Negator ( inChannel: c.in(),
                                     outChannel: d.out() ),
                             new GPlus ( inChannel0: a.in(),
                                         inChannel1: d.in(),
                                         outChannel: outChannel )
                     ]

   new PAR ( differentiateList ).run()

 }
```
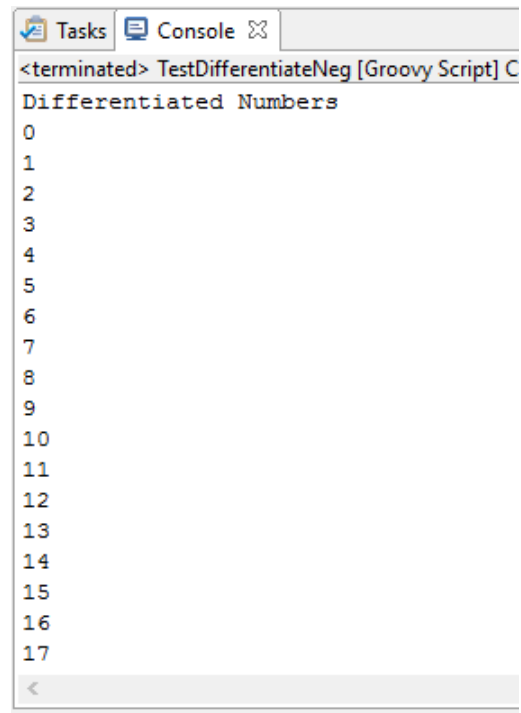
3.3   Minus

```
void run () {

  ProcessRead read0 = new ProcessRead ( inChannel0)
  ProcessRead read1 = new ProcessRead ( inChannel1)
  def parRead2 = new PAR ( [ read0, read1 ] )

  while (true) {
    parRead2.run()
    // output one value subtracted from the other
    // be certain you know which way round you are doing the
        subtraction!!
   outChannel.write(read0.value - read1.value)

  }
}
```

3.4   Negator

```
void run () {
  while (true) {
    //output the negative of the input value
    outChannel.write(-inChannel.read())
  }
}
```

Figure 10: Results from running DifferentiateNegator.

4 Exercise 3.2

Q: Determine the effect of the change and why it happens.

Ans: The output channels of pairs A and B are switched. With GSPairsA if an output isn't ready to receive data the process deadlocks as GSPairsA writes to GPlus first and GSPairsB doesn't.
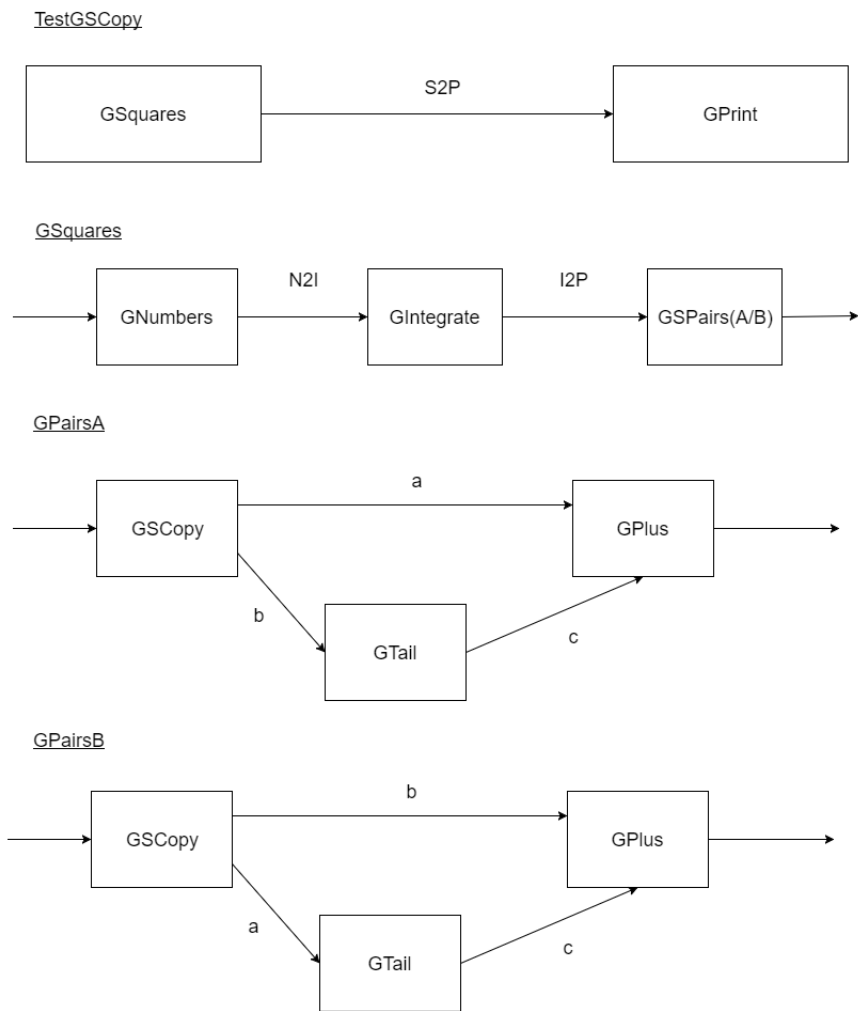
TestGSCopy

```
┌─────────────────┐                    ┌─────────────────┐
│                 │        S2P         │                 │
│    GSquares     │ ─────────────────▶ │     GPrint      │
│                 │                    │                 │
└─────────────────┘                    └─────────────────┘
```

GSquares

```
        ┌───────────┐   N2I   ┌───────────┐   I2P   ┌───────────────┐
──────▶ │ GNumbers  │ ──────▶ │ GIntegrate│ ──────▶ │ GSPairs(A/B)  │ ──────▶
        └───────────┘         └───────────┘         └───────────────┘
```

GPairsA

```
        ┌───────────┐            a              ┌───────────┐
──────▶ │  GSCopy   │ ────────────────────────▶ │   GPlus   │ ──────▶
        └───────────┘                           └───────────┘
              │  b                                    ▲
              ▼              ┌───────────┐      c     │
                            │   GTail   │ ───────────┘
                            └───────────┘
```

GPairsB

```
        ┌───────────┐            b              ┌───────────┐
──────▶ │  GSCopy   │ ────────────────────────▶ │   GPlus   │ ──────▶
        └───────────┘                           └───────────┘
              │  a                                    ▲
              ▼              ┌───────────┐      c     │
                            │   GTail   │ ───────────┘
                            └───────────┘
```

Figure 11: Process Diagrams for exercise 3-2.
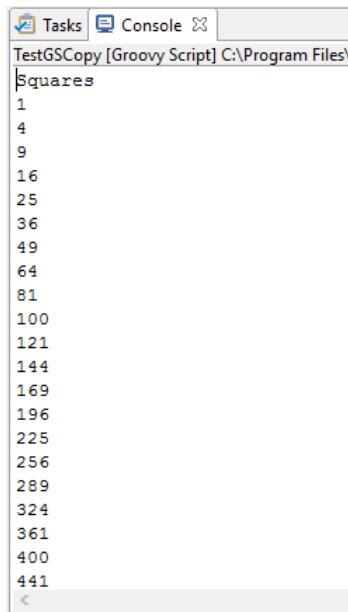
Figure 12: TestGSCopy program running.

4.1 GSCopy

```groovy
void run () {

    def write0 = new ProcessWrite(outChannel0)
    def write1 = new ProcessWrite(outChannel1)
    def parWrite2 = new PAR([write0, write1])

    while (true) {
     def i = inChannel.read()
     // output the input value in sequence to each output channel
    write0.value = i
    write1.value = i
    parWrite2.run()

  }
}
```

4.2 GSquares

```
void run () {

  One2OneChannel N2I = Channel.createOne2One()
  One2OneChannel I2P = Channel.createOne2One()

  def testList = [ new GNumbers ( outChannel: N2I.out() ),
                   new GIntegrate ( inChannel: N2I.in(),
                                    outChannel: I2P.out() ),
              new GSPairsB ( inChannel: I2P.in(),
                             outChannel: outChannel )
                  // you will need to modify this twice
                  //first modification is to insert a constructor for
                      GSPairsA
                  // then run the network using TestGSCopy
                  //second modification replace the constructor for
                      GSPairsA with GSPairsB
                  // then run the network again using TestGSCopy
                  // you will then be able to compare the behaviour
                      and also to
                  // explain why this happens!
              //GSPairsA does not work
                 ]
  new PAR ( testList ).run()
}
```

5   Exercise 3.3

Q: Why was it considered easier to build GParPrint as a new process rather than using multiple instances of GPrint to output the table of results?

Ans: GParPrint makes for easier processing using a collection of channels. This is done using the GPCopy process at intermediate stages in the pipeline depending on what data is to be shown in the table whilst still passing on the data to the next process. Unlike GPrint, GParPrint prints the data in a table-like format with columns and makes the data more readable where GPrint prints the data in rows and makes it less readable due to the complex routing of the channels. Doing this also reduces the chances of errors by simplifying the system.

6   Exercise 4.1

Q: What happens if line 25 of ResetPrefix Listing 4-1 is commented out? Why?

Ans: The Alt can switch which channel it recieves its input from, the resetChan-

nel and the inChannel. In this case the resetChannel is given priority over the inChannel. An initial number is set and this number is going upwards in increments of 1 and sent to the outChannel. If there is any input from the user it will come in through the resetChannel and the program should begin incrementing by 1 from this new value. Once the new input data has been read the system must wait for the outputChannel to be ready to receive data from the inChannel. This is verified by waiting for the inChannel.read (line 25) to return a value.

If line 25 is commented out the program will continue to increment the previous values whilst also outputting the new values which are being incremented by 1, if another input is entered then the system deadlocks. This happens due to the lack of processes which can take in new numbers, this could be overcome by possibly adding a buffer or simply more processes.
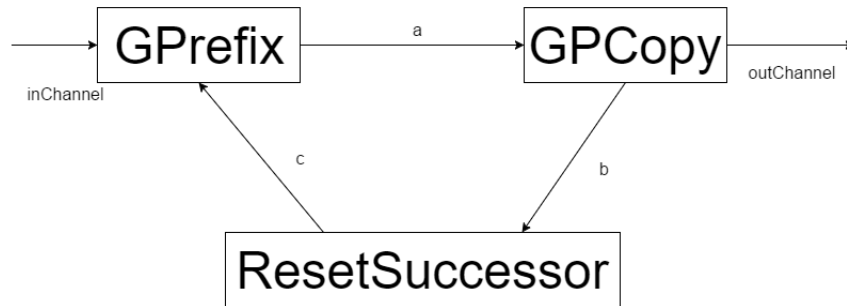
Figure 13: Diagram of the ResetNumbers process.

7 Exercise 4.2

Q: You will have to write a ResetSuccessor process. Does it overcome the problem identified in the previous exercise? If not, why not?

Ans: No this does not overcome the problem identified. This is because the ResetSuccessor process needs to read from the last iteration before continuing any further in the cycle.
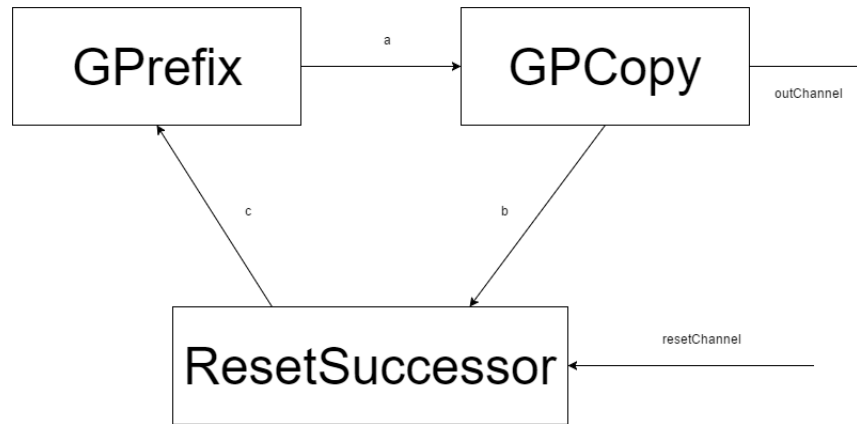
Figure 14: Diagram of the ResetNumbers process.

7.1 ResetSuccessor

```
void run () {
  def guards = [ resetChannel, inChannel ]
  def alt = new ALT ( guards )
  while (true) {
   // deal with inputs from resetChannel and inChannel
   // use a priSelect

    def index = alt.priSelect()
    if(index == 0) { //Reset channel input
       def resetValue = resetChannel.read()
       inChannel.read()
       outChannel.write(resetValue)
    }
    else {  //inChannel input
       outChannel.write(inChannel.read()+1)
    }

 }
 }
```

7.2 ResetNumbers

```
void run() {

  One2OneChannel a = Channel.createOne2One()
```

```
    One2OneChannel b = Channel.createOne2One()
    One2OneChannel c = Channel.createOne2One()

    def testList = [ new GPrefix ( prefixValue: initialValue,
                                    outChannel: a.out(),
                                    inChannel: c.in() ),
                  new GPCopy ( inChannel: a.in(),
                               outChannel0: outChannel,
                               outChannel1: b.out() ),
                  // requires a constructor for ResetSuccessor
              new ResetSuccessor ( inChannel: b.in(),
                                   outChannel: c.out(),
                                    resetChannel: resetChannel )
                ]
  new PAR ( testList ).run()
}
```
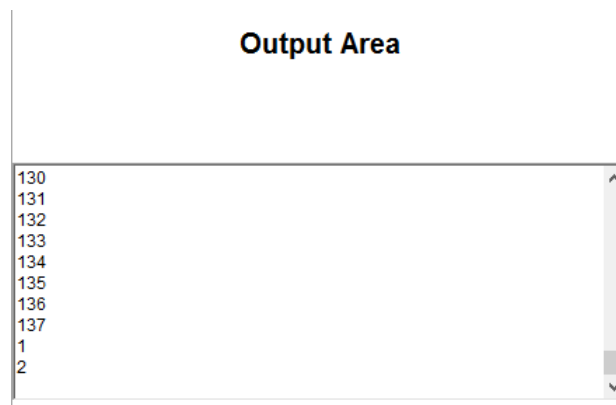

Figure 15: System output with inChannel.read()
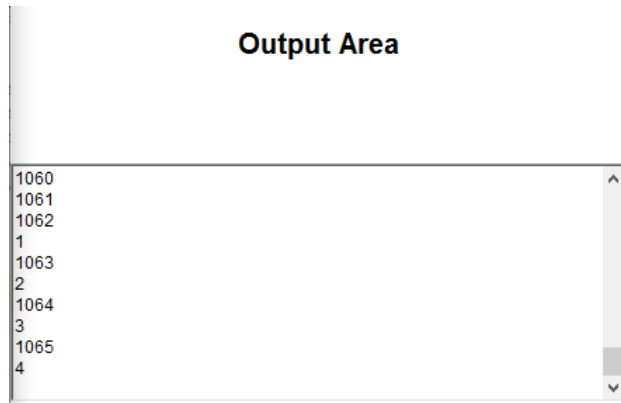
Figure 16: System output without inChannel.read() results in deadlock

8 Exercise 5.1

Q: What do you conclude from these experiments?

Ans: These results show that the delay effects the program equally when it is imposed on either of the processes. This happens because the two processes loop until completion and the delays begin at the same time rather than one after the other so after the initial delay the following process is ready to take in the data and process it.
This also shows that the Producer writes data to the queue based on the delay timer and doesn't need to know the state of the Consumer.

Table 1: Table of results.

Queue Delay Results		
QProducer Delay (ms)	Queue Delay (ms)	Total Time (secs)
0	0	1.85
500	0	26.77
1000	0	51.53
0	500	25.63
0	1000	47.49
500	500	26.69
1000	500	51.77
500	1000	51.53
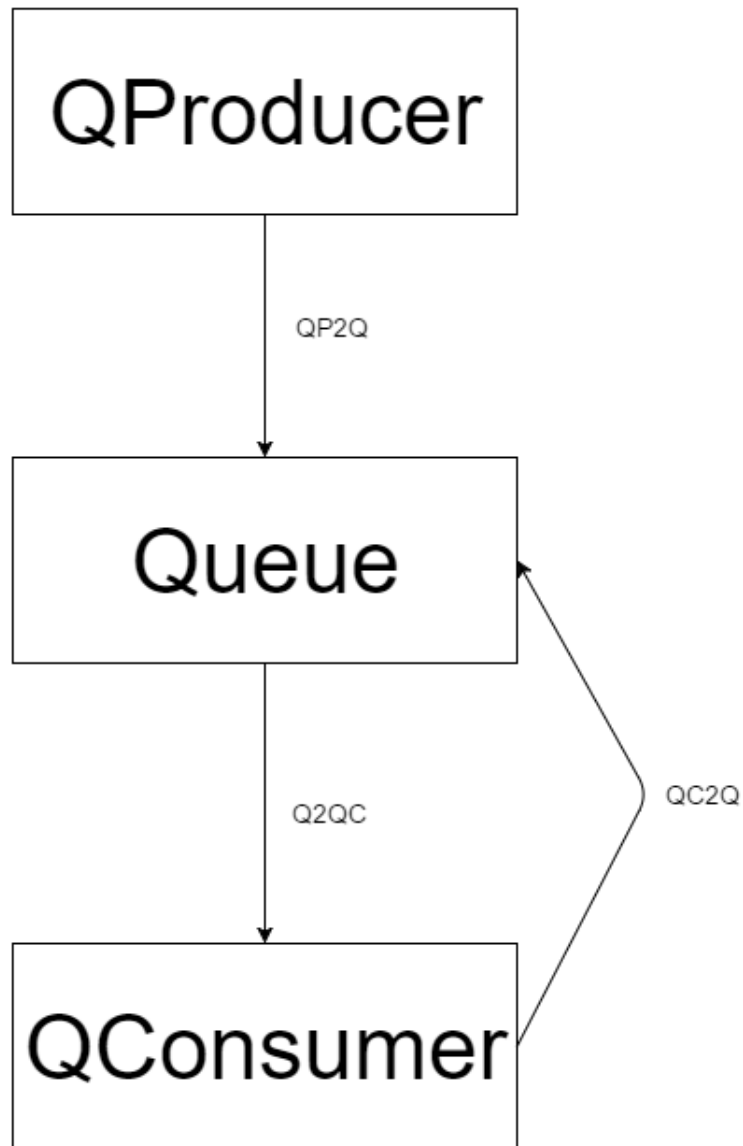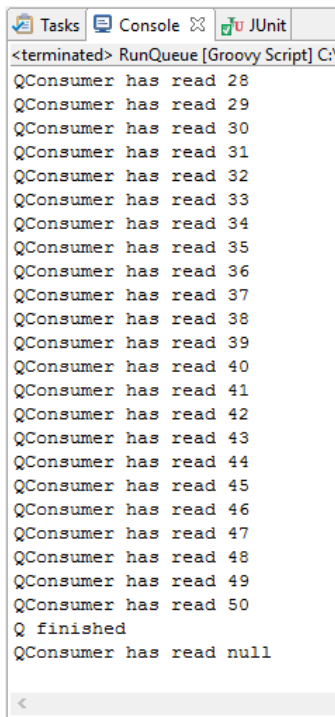
Figure 17: Queue system.

Figure 18: Queue system.

9 Exercise 5.2

Q: Which is the more elegant formulation? Why?

Ans: This solution is more elegant. The code is compact and far clearer/easier to read. This is because only one Alt and one switch statement is used in this solution. States caneasily be switched back and forth.

9.1 Scale

```
while (true) {
    switch ( scaleAlt.priSelect(preCon) ) {
      case SUSPEND :
        // deal with suspend input
    suspend.read()
    factor.write(scaling)
```

```
        println "SUSPEND"
        preCon[INJECT] = true
        preCon[SUSPEND] = true
            break
          case INJECT:
            // deal with inject input
        scaling = injector.read()
        println "INJECTED SCALING: $scaling"
        timeout = timer.read() + DOUBLE_INTERVAL
        timer.setAlarm(timeout)
        preCon[SUSPEND] = true
        preCon[INJECT] = false
            break
          case TIMER:
            // deal with Timer input
        timeout = timer.read() + DOUBLE_INTERVAL
        timer.setAlarm(timeout)
        scaling = scaling * 2
            println "NORMAL TIMER: NEW SCALING: ${scaling}"
            break
          case INPUT:
            //  deal with Input channel
        def inValue = inChannel.read()
        def result = new ScaledData()
        result.original = inValue
        if(preCon[SUSPEND]){
            result.scaled = inValue * scaling
        }
        else{
            result.scaled = inValue
        }
        outChannel.write(result)
            break
        } //end-switch
    } //end-while
```

Figure 19: Scaling System output.

10 Exercise 6.1

10.1 TestCase

```groovy
class TestCase extends GroovyTestCase {

  void test() {

  One2OneChannel connect1 = Channel.createOne2One()
  One2OneChannel connect2 = Channel.createOne2One()

  def generate = new GenerateSetsOfThree(outChannel: connect1.out())

  def List2Strm = new ListToStream(inChannel: connect1.in(),
                       outChannel: connect2.out())
```

```
    def CreateSetsOEight = new CreateSetsOfEight(inChannel: connect2.in())

    def testList = [generate, List2Strm, CreateSetsOEight]
    new PAR (testList).run()

    def expected = generate.sequence
    def actual = CreateSetsOEight.OutSeq
    assertTrue(expected == actual)


    }

}
```

10.2   CreateSetsOfEight

```
class CreateSetsOfEight implements CSProcess{

    def ChannelInput inChannel
    def int listParam = 7

    def OutSeq = []

    void run(){
        def outList = []
        def v = inChannel.read()
        while (v != -1){
            for ( i in 0 .. listParam ) {
                // put v into outList and read next input
                outList << v

                OutSeq = OutSeq << v

                v = inChannel.read()
            }
            println " Eight Object is ${outList}"
            outList = []
        }
        println "Finished"
    }
}
```

10.3   GenerateSetsOfThree

```
class GenerateSetsOfThree implements CSProcess {
```

```
def ChannelOutput outChannel

def sequence = []

void run(){
  def threeList = [
                  [1, 2, 3],
                  [4, 5, 6],
                  [7, 8, 9],
                  [10, 11, 12],
                  [13, 14, 15],
                  [16, 17, 18],
                  [19, 20, 21],
                  [22, 23, 24] ]
  for ( i in 0 ..< threeList.size)
  {
    outChannel.write(threeList[i])

    sequence = sequence << threeList[i]
  }
  //write the terminating List as per exercise definition
  def terminatingList = [-1, -1, -1]
  outChannel.write(terminatingList)

  sequence = sequence.flatten()
  println "TERMINATING: $terminatingList $sequence "
  }
}
```
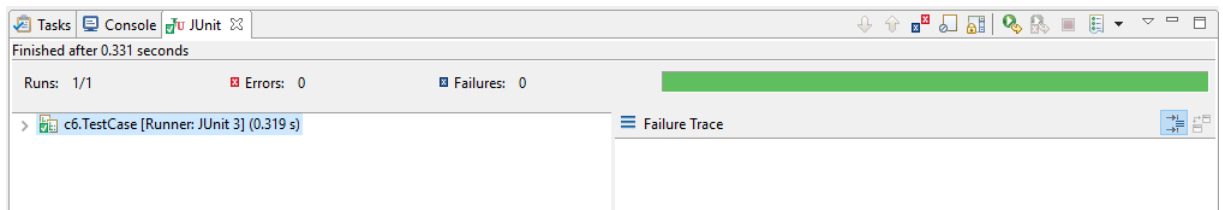

Figure 20: Test case working 6-1.

11   Exercise 7.1

Q: Determine the cause of the deadlock.

Ans: The deadlock only happens when both servers are communicating with each other at the same time and are both requesting data. They are both

waiting to receive the data and stuck in a continuous loop forever causing the deadlock.
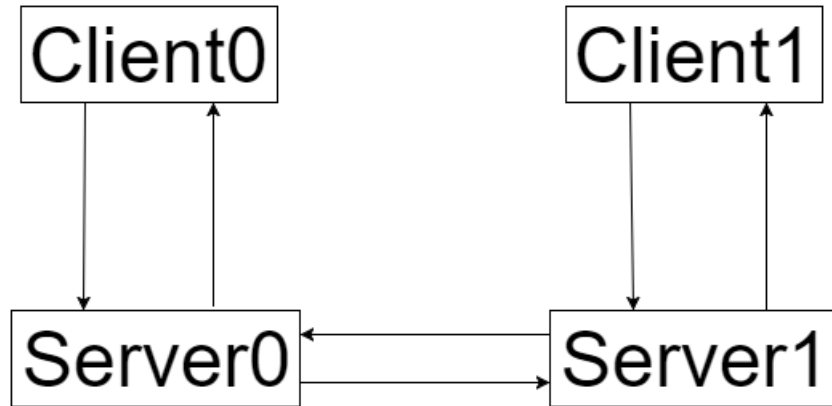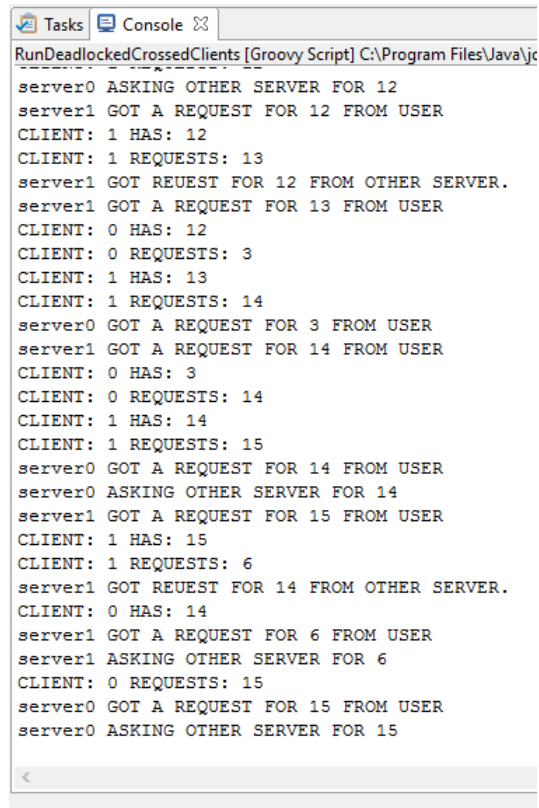
Figure 21: System diagram for exercise 7-1.

Figure 22: System output for exercise 7-1.

12 Exercise 8.1

12.1 Client

```groovy
void run () {
   def iterations = selectList.size
   println "CLIENT: $clientNumber HAS: $iterations VALUES IN:
       $selectList"

   for ( i in 0 ..< iterations) {
     def key = selectList[i]
    println "CLIENT: $clientNumber REQUESTS: $key"
     requestChannel.write(key)
     def v = receiveChannel.read()
    println "CLIENT: $clientNumber HAS: $key"
```

28

```
def testVal = key*10

if(testVal == v) {
   println("KEY $clientNumber READ $v")
   println "TEST SUCCESS"
}
else {
   println("KEY $clientNumber READ $v")
   println "TEST FAILED"
}
}

println "CLIENT: $clientNumber FINISHED."
}
```


Figure 23: Output for exercise 8-1.

13 Exercise 9.1

13.1 MissCheck

```
class MissCheck implements CSProcess {

  def ChannelInput inChannel
  def ChannelOutput outChannel
  int preVal = 99

  void run() {
    while (true)
    {
      def inStream = inChannel.read()
      outChannel.write(inStream)

      def missed = ( inStream.data - preVal) - 1
      if(missed != inStream.missed ){
        println "EVENT W/ DATA: $inStream.data MISSED:
            $inStream.missed CORRECTION: $missed"
      }
      preVal = inStream.data
    }
  }
}
```

---

14 Exercise 9.2

Q: What do you conclude?

Ans: When the program is first run, for the most part the first few events pass through without being missed. As seen in the figure below (Fig:24) during the first few cycles the limited buffers will be empty. If the buffers were increased in size this would reduce the amount of missed events.
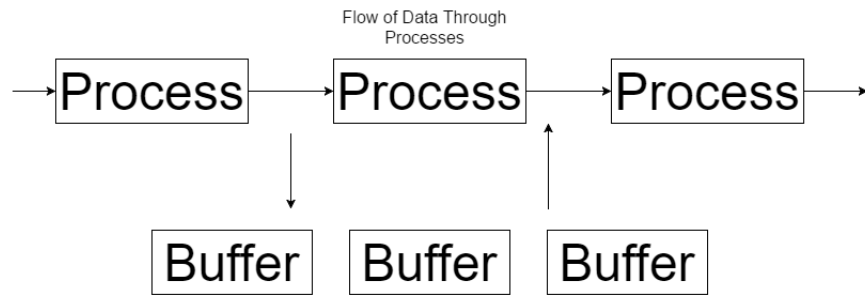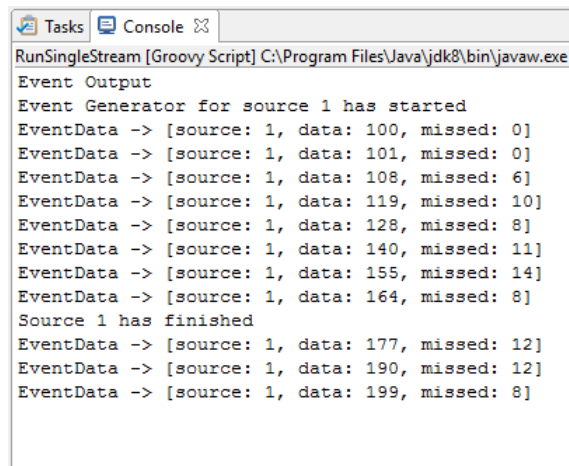
Figure 24: An abstract diagram to help illustrate the point.

15 Exercise 9.3

Fairmultiplex: This method chooses a channel based on what priority it is set to. It works similarly to a queuing system, after a channel has been selected it is then allocated a low priority during the next selection process. This gives consistent output.
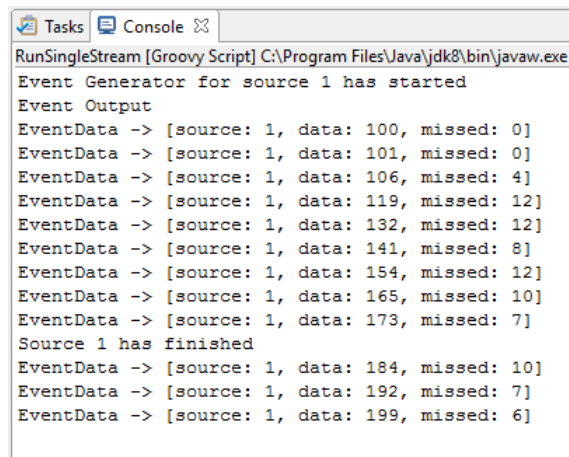
Primultiplex: This method chooses a channel based on which channel is ready first. This results in the prioritization of streams with the a lower index i.e stream 1. This may result in any events in streams with a higher index not being read.

Muliplexer: This method randomly chooses between the different channels that are currently ready. As a result this method produce the most inconsistent output.

Figure 25: Results using Fairmultiplex.

Figure 26: Results using Primultiplex.

Figure 27: Results using Multiplexer.

16 Exercise 11.3

17 Requirement 2

After looking at the code it is apparent there is an abundance of nested while loops and if statements, this makes the code relatively clunky and could be made more readable. When looking at how the system is structured and after running it, it is apparent that the system actually deadlocks. The deadlock appears to occur between the matcher and mouse buffer and this problem could probably be solved by introducing a queue process which acts as an intermediate buffer and would regulate the flow of data.